

```

struct task_struct {
#ifdef CONFIG_THREAD_INFO_IN_TASK
    /*
     * For reasons of header soup (see current_thread_info()), this
     * must be the first element of task_struct.
     */
    struct thread_info      thread_info;
#endif

    /* -1 unrunnable, 0 runnable, >0 stopped: */
    volatile long          state;

    /*
     * This begins the randomizable portion of task_struct. Only
     * scheduling-critical items should be added above here.
     */
    randomized_struct_fields_start

    void                  *stack;
    refcount_t            usage;
    /* Per task flags (PF_*), defined further below: */
    unsigned int          flags;
    unsigned int          ptrace;

#ifdef CONFIG_SMP
    struct llist_node      wake_entry;
    int                    on_cpu;
#endif
#ifdef CONFIG_THREAD_INFO_IN_TASK
    /* Current CPU: */
    unsigned int          cpu;
#endif
    unsigned int          wakee_flips;
    unsigned long         wakee_flip_decay_ts;
    struct task_struct     *last_wakee;

    /*
     * recent_used_cpu is initially set as the last CPU used by a task
     * that wakes affine another task. Waker/wakee relationships can
     * push tasks around a CPU where each wakeup moves to the next one.
     * Tracking a recently used CPU allows a quick search for a recently
     * used CPU that may be idle.
     */
    int                    recent_used_cpu;
    int                    wake_cpu;
#endif
    int                    on_rq;

```

```

int                                prio;
int                                static_prio;
int                                normal_prio;
unsigned int                        rt_priority;

const struct sched_class            *sched_class;
struct sched_entity                se;
struct sched_rt_entity              rt;
#ifdef CONFIG_CGROUP_SCHED
struct task_group                    *sched_task_group;
#endif
struct sched_dl_entity              dl;

#ifdef CONFIG_UCLAMP_TASK
/* Clamp values requested for a scheduling entity */
struct uclamp_se                    uclamp_req[UCLAMP_CNT];
/* Effective clamp values used for a scheduling entity */
struct uclamp_se                    uclamp[UCLAMP_CNT];
#endif

#ifdef CONFIG_PREEMPT_NOTIFIERS
/* List of struct preempt_notifier: */
struct hlist_head                    preempt_notifiers;
#endif

#ifdef CONFIG_BLK_DEV_IO_TRACE
unsigned int                        btrace_seq;
#endif

unsigned int                        policy;
int                                nr_cpus_allowed;
const cpumask_t                    *cpus_ptr;
cpumask_t                          cpus_mask;

#ifdef CONFIG_PREEMPT_RCU
int                                rcu_read_lock_nesting;
union rcu_special                    rcu_read_unlock_special;
struct list_head                    rcu_node_entry;
struct rcu_node                      *rcu_blocked_node;
#endif /* #ifdef CONFIG_PREEMPT_RCU */

#ifdef CONFIG_TASKS_RCU
unsigned long                        rcu_tasks_nvcsw;
u8                                  rcu_tasks_holdout;
u8                                  rcu_tasks_idx;
int                                rcu_tasks_idle_cpu;

```

```

        struct list_head          rcu_tasks_holdout_list;
#endif /* #ifdef CONFIG_TASKS_RCU */

        struct sched_info         sched_info;

        struct list_head          tasks;
#ifdef CONFIG_SMP
        struct plist_node         pushable_tasks;
        struct rb_node            pushable_dl_tasks;
#endif

        struct mm_struct          *mm;
        struct mm_struct          *active_mm;

        /* Per-thread vma caching: */
        struct vmacache           vmacache;

#ifdef SPLIT_RSS_COUNTING
        struct task_rss_stat       rss_stat;
#endif

        int                       exit_state;
        int                       exit_code;
        int                       exit_signal;
        /* The signal sent when the parent dies: */
        int                       pdeath_signal;
        /* JOBCTL_*, siglock protected: */
        unsigned long             jobctl;

        /* Used for emulating ABI behavior of previous Linux versions: */
        unsigned int              personality;

        /* Scheduler bits, serialized by scheduler locks: */
        unsigned                   sched_reset_on_fork:1;
        unsigned                   sched_contributes_to_load:1;
        unsigned                   sched_migrated:1;
        unsigned                   sched_remote_wakeup:1;
#ifdef CONFIG_PSI
        unsigned                   sched_psi_wake_requeue:1;
#endif

        /* Force alignment to the next boundary: */
        unsigned                   :0;

        /* Unserialized, strictly 'current' */

        /* Bit to tell LSMs we're in execve(): */

```

```

        unsigned                in_execve:1;
        unsigned                in_iowait:1;
#ifdef TIF_RESTORE_SIGMASK
        unsigned                restore_sigmask:1;
#endif
#ifdef CONFIG_MEMCG
        unsigned                in_user_fault:1;
#endif
#ifdef CONFIG_COMPAT_BRK
        unsigned                brk_randomized:1;
#endif
#ifdef CONFIG_CGROUPS
        /* disallow userland-initiated cgroup migration */
        unsigned                no_cgroup_migration:1;
        /* task is frozen/stopped (used by the cgroup freezer) */
        unsigned                frozen:1;
#endif
#ifdef CONFIG_BLK_CGROUP
        /* to be used once the psi infrastructure lands upstream. */
        unsigned                use_memdelay:1;
#endif

        unsigned long           atomic_flags; /* Flags requiring atomic access. */

        struct restart_block     restart_block;

        pid_t                   pid;
        pid_t                   tgid;

#ifdef CONFIG_STACKPROTECTOR
        /* Canary value for the -fstack-protector GCC feature: */
        unsigned long           stack_canary;
#endif

    /*
     * Pointers to the (original) parent process, youngest child, younger sibling,
     * older sibling, respectively. (p->father can be replaced with
     * p->real_parent->pid)
     */

    /* Real parent process: */
    struct task_struct __rcu     *real_parent;

    /* Recipient of SIGCHLD, wait4() reports: */
    struct task_struct __rcu     *parent;

    /*

```

```

    * Children/sibling form the list of natural children:
    */
    struct list_head      children;
    struct list_head      sibling;
    struct task_struct    *group_leader;

    /*
     * 'ptraced' is the list of tasks this task is using ptrace() on.
     *
     * This includes both natural children and PTRACE_ATTACH targets.
     * 'ptrace_entry' is this task's link on the p->parent->ptraced list.
     */
    struct list_head      ptraced;
    struct list_head      ptrace_entry;

    /* PID/PID hash table linkage. */
    struct pid            *thread_pid;
    struct hlist_node     pid_links[PIDTYPE_MAX];
    struct list_head      thread_group;
    struct list_head      thread_node;

    struct completion     *vfork_done;

    /* CLONE_CHILD_SETTID: */
    int __user            *set_child_tid;

    /* CLONE_CHILD_CLEARTID: */
    int __user            *clear_child_tid;

    u64                   utime;
    u64                   stime;
#ifdef CONFIG_ARCH_HAS_SCALED_CPUTIME
    u64                   utimescaled;
    u64                   stimescaled;
#endif
    u64                   gtime;
    struct prev_cputime    prev_cputime;
#ifdef CONFIG_VIRT_CPU_ACCOUNTING_GEN
    struct vtime           vtime;
#endif

#ifdef CONFIG_NO_HZ_FULL
    atomic_t              tick_dep_mask;
#endif

    /* Context switch counts: */
    unsigned long         nvcs;

```

```

unsigned long                nivcsw;

/* Monotonic time in nsecs: */
u64                          start_time;

/* Boot based time in nsecs: */
u64                          start_boottime;

/* MM fault and swap info: this can arguably be seen as either mm-specific or
thread-specific: */
unsigned long                minflt;
unsigned long                majflt;

/* Empty if CONFIG_POSIX_CPUTIMERS=n */
struct posix_cputimers       posix_cputimers;

/* Process credentials: */

/* Tracer's credentials at attach: */
const struct cred __rcu      *ptracer_cred;

/* Objective and real subjective task credentials (COW): */
const struct cred __rcu      *real_cred;

/* Effective (overridable) subjective task credentials (COW): */
const struct cred __rcu      *cred;

#ifdef CONFIG_KEYS
/* Cached requested key. */
struct key                   *cached_requested_key;
#endif

/*
 * executable name, excluding path.
 *
 * - normally initialized setup_new_exec()
 * - access it with [gs]et_task_comm()
 * - lock it with task_lock()
 */
char                          comm[TASK_COMM_LEN];

struct nameidata              *nameidata;

#ifdef CONFIG_SYSVIPC
struct sysv_sem               sysvsem;
struct sysv_shm               sysvshm;

```

```

#endif
#ifdef CONFIG_DETECT_HUNG_TASK
    unsigned long        last_switch_count;
    unsigned long        last_switch_time;
#endif

    /* Filesystem information: */
    struct fs_struct      *fs;

    /* Open file information: */
    struct files_struct   *files;

    /* Namespaces: */
    struct nsproxy        *nsproxy;

    /* Signal handlers: */
    struct signal_struct  *signal;
    struct sighand_struct *sighand;
    sigset_t              blocked;
    sigset_t              real_blocked;
    /* Restored if set_restore_sigmask() was used: */
    sigset_t              saved_sigmask;
    struct sigpending     pending;
    unsigned long         sas_ss_sp;
    size_t                sas_ss_size;
    unsigned int          sas_ss_flags;

    struct callback_head  *task_works;

#ifdef CONFIG_AUDIT
#ifdef CONFIG_AUDITSYSCALL
    struct audit_context  *audit_context;
#endif
#endif

    kuid_t                loginuid;
    unsigned int           sessionid;
#endif

    struct seccomp         seccomp;

    /* Thread group tracking: */
    u32                    parent_exec_id;
    u32                    self_exec_id;

    /* Protection against (de-)allocation: mm, files, fs, tty, keyrings, mems_allowed,
mempolicy: */
    spinlock_t             alloc_lock;

    /* Protection of the PI data structures: */

```

```

raw_spinlock_t                pi_lock;

struct wake_q_node            wake_q;

#ifdef CONFIG_RT_MUTEXES
    /* PI waiters blocked on a rt_mutex held by this task: */
    struct rb_root_cached      pi_waiters;
    /* Updated under owner's pi_lock and rq lock */
    struct task_struct          *pi_top_task;
    /* Deadlock detection and priority inheritance handling: */
    struct rt_mutex_waiter      *pi_blocked_on;
#endif

#ifdef CONFIG_DEBUG_MUTEXES
    /* Mutex deadlock detection: */
    struct mutex_waiter         *blocked_on;
#endif

#ifdef CONFIG_DEBUG_ATOMIC_SLEEP
    int                         non_block_count;
#endif

#ifdef CONFIG_TRACE_IRQFLAGS
    unsigned int                irq_events;
    unsigned long               hardirq_enable_ip;
    unsigned long               hardirq_disable_ip;
    unsigned int                hardirq_enable_event;
    unsigned int                hardirq_disable_event;
    int                         hardirqs_enabled;
    int                         hardirq_context;
    unsigned long               softirq_disable_ip;
    unsigned long               softirq_enable_ip;
    unsigned int                softirq_disable_event;
    unsigned int                softirq_enable_event;
    int                         softirqs_enabled;
    int                         softirq_context;
#endif

#ifdef CONFIG_LOCKDEP
    #define MAX_LOCK_DEPTH      48UL
    u64                         curr_chain_key;
    int                         lockdep_depth;
    unsigned int                lockdep_recursion;
    struct held_lock             held_locks[MAX_LOCK_DEPTH];
#endif

```



```

#ifdef CONFIG_UBSAN
    unsigned int                in_ubsan;
#endif

    /* Journalling filesystem info: */
    void                        *journal_info;

    /* Stacked block device info: */
    struct bio_list              *bio_list;

#ifdef CONFIG_BLOCK
    /* Stack plugging: */
    struct blk_plug              *plug;
#endif

    /* VM state: */
    struct reclaim_state         *reclaim_state;

    struct backing_dev_info      *backing_dev_info;

    struct io_context            *io_context;

#ifdef CONFIG_COMPACTION
    struct capture_control       *capture_control;
#endif

    /* Ptrace state: */
    unsigned long                ptrace_message;
    kernel_siginfo_t             *last_siginfo;

    struct task_io_accounting     ioac;
#ifdef CONFIG_PSI
    /* Pressure stall state */
    unsigned int                 psi_flags;
#endif
#ifdef CONFIG_TASK_XACCT
    /* Accumulated RSS usage: */
    u64                          acct_rss_mem1;
    /* Accumulated virtual memory usage: */
    u64                          acct_vm_mem1;
    /* stime + utime since last update: */
    u64                          acct_timexpd;
#endif
#ifdef CONFIG_CPUSETS
    /* Protected by ->alloc_lock: */
    nodemask_t                   mems_allowed;
    /* Sequence number to catch updates: */

```

```

        seqcount_t                mems_allowed_seq;
        int                       cpuset_mem_spread_rotor;
        int                       cpuset_slab_spread_rotor;
#endif
#ifdef CONFIG_CGROUPS
    /* Control Group info protected by css_set_lock: */
    struct css_set __rcu          *cgroups;
    /* cg_list protected by css_set_lock and task->alloc_lock: */
    struct list_head              cg_list;
#endif
#ifdef CONFIG_X86_CPU_RESCTRL
    u32                           closid;
    u32                           rmid;
#endif
#ifdef CONFIG_FUTEX
    struct robust_list_head __user *robust_list;
#endif
#ifdef CONFIG_COMPAT
    struct compat_robust_list_head __user *compat_robust_list;
#endif
    struct list_head              pi_state_list;
    struct futex_pi_state         *pi_state_cache;
    struct mutex                  futex_exit_mutex;
    unsigned int                  futex_state;
#endif
#ifdef CONFIG_PERF_EVENTS
    struct perf_event_context     *perf_event_ctxp[perf_nr_task_contexts];
    struct mutex                  perf_event_mutex;
    struct list_head              perf_event_list;
#endif
#ifdef CONFIG_DEBUG_PREEMPT
    unsigned long                 preempt_disable_ip;
#endif
#ifdef CONFIG_NUMA
    /* Protected by alloc_lock: */
    struct mempolicy              *mempolicy;
    short                         il_prev;
    short                         pref_node_fork;
#endif
#ifdef CONFIG_NUMA_BALANCING
    int                           numa_scan_seq;
    unsigned int                  numa_scan_period;
    unsigned int                  numa_scan_period_max;
    int                           numa_preferred_nid;
    unsigned long                 numa_migrate_retry;
    /* Migration stamp: */
    u64                           node_stamp;

```

```

u64                                last_task_numa_placement;
u64                                last_sum_exec_runtime;
struct callback_head               numa_work;

/*
 * This pointer is only modified for current in syscall and
 * pagefault context (and for tasks being destroyed), so it can be read
 * from any of the following contexts:
 * - RCU read-side critical section
 * - current->numa_group from everywhere
 * - task's runqueue locked, task not running
 */
struct numa_group __rcu            *numa_group;

/*
 * numa_faults is an array split into four regions:
 * faults_memory, faults_cpu, faults_memory_buffer, faults_cpu_buffer
 * in this precise order.
 *
 * faults_memory: Exponential decaying average of faults on a per-node
 * basis. Scheduling placement decisions are made based on these
 * counts. The values remain static for the duration of a PTE scan.
 * faults_cpu: Track the nodes the process was running on when a NUMA
 * hinting fault was incurred.
 * faults_memory_buffer and faults_cpu_buffer: Record faults per node
 * during the current scan window. When the scan completes, the counts
 * in faults_memory and faults_cpu decay and these values are copied.
 */
unsigned long                      *numa_faults;
unsigned long                      total_numa_faults;

/*
 * numa_faults_locality tracks if faults recorded during the last
 * scan window were remote/local or failed to migrate. The task scan
 * period is adapted based on the locality of the faults with different
 * weights depending on whether they were shared or private faults
 */
unsigned long                      numa_faults_locality[3];

unsigned long                      numa_pages_migrated;
#endif /* CONFIG_NUMA_BALANCING */

#ifdef CONFIG_RSEQ
struct rseq __user *rseq;
u32 rseq_sig;
/*

```

```

    * RmW on rseq_event_mask must be performed atomically
    * with respect to preemption.
    */
    unsigned long rseq_event_mask;
#endif

    struct tlbflush_unmap_batch tlb_ubic;

    union {
        refcount_t          rcu_users;
        struct rcu_head      rcu;
    };

    /* Cache last used pipe for splice(): */
    struct pipe_inode_info   *splice_pipe;

    struct page_frag        task_frag;

#ifdef CONFIG_TASK_DELAY_ACCT
    struct task_delay_info   *delays;
#endif

#ifdef CONFIG_FAULT_INJECTION
    int                      make_it_fail;
    unsigned int             fail_nth;
#endif

    /*
     * When (nr_dirtied >= nr_dirtied_pause), it's time to call
     * balance_dirty_pages() for a dirty throttling pause:
     */
    int                      nr_dirtied;
    int                      nr_dirtied_pause;
    /* Start of a write-and-pause period: */
    unsigned long            dirty_paused_when;

#ifdef CONFIG_LATENCYTOP
    int                      latency_record_count;
    struct latency_record     latency_record[LT_SAVECOUNT];
#endif

    /*
     * Time slack values; these are used to round up poll() and
     * select() etc timeout values. These are in nanoseconds.
     */
    u64                      timer_slack_ns;
    u64                      default_timer_slack_ns;

```

```

#ifdef CONFIG_KASAN
    unsigned int            kasan_depth;
#endif

#ifdef CONFIG_FUNCTION_GRAPH_TRACER
    /* Index of current stored address in ret_stack: */
    int                    curr_ret_stack;
    int                    curr_ret_depth;

    /* Stack of return addresses for return function tracing: */
    struct ftrace_ret_stack *ret_stack;

    /* Timestamp for last schedule: */
    unsigned long long     ftrace_timestamp;

    /*
     * Number of functions that haven't been traced
     * because of depth overrun:
     */
    atomic_t               trace_overrun;

    /* Pause tracing: */
    atomic_t               tracing_graph_pause;
#endif

#ifdef CONFIG_TRACING
    /* State flags for use by tracers: */
    unsigned long          trace;

    /* Bitmask and counter of trace recursion: */
    unsigned long          trace_recursion;
#endif /* CONFIG_TRACING */

#ifdef CONFIG_KCOV
    /* See kernel/kcov.c for more details. */

    /* Coverage collection mode enabled for this task (0 if disabled): */
    unsigned int           kcov_mode;

    /* Size of the kcov_area: */
    unsigned int           kcov_size;

    /* Buffer for coverage collection: */
    void                   *kcov_area;

    /* KCOV descriptor wired with this task or NULL: */

```

```

        struct kcov                *kcov;

        /* KCOV common handle for remote coverage collection: */
        u64                        kcov_handle;

        /* KCOV sequence number: */
        int                        kcov_sequence;
#endif

#ifdef CONFIG_MEMCG
        struct mem_cgroup          *memcg_in_oom;
        gfp_t                      memcg_oom_gfp_mask;
        int                        memcg_oom_order;

        /* Number of pages to reclaim on returning to userland: */
        unsigned int              memcg_nr_pages_over_high;

        /* Used by memcontrol for targeted memcg charge: */
        struct mem_cgroup          *active_memcg;
#endif

#ifdef CONFIG_BLK_CGROUP
        struct request_queue       *throttle_queue;
#endif

#ifdef CONFIG_UPROBES
        struct uprobe_task        *utask;
#endif

#ifdef CONFIG_BCACHE || defined(CONFIG_BCACHE_MODULE)
        unsigned int              sequential_io;
        unsigned int              sequential_io_avg;
#endif

#ifdef CONFIG_DEBUG_ATOMIC_SLEEP
        unsigned long             task_state_change;
#endif

        int                        pagefault_disabled;
#ifdef CONFIG_MMU
        struct task_struct         *oom_reaper_list;
#endif

#ifdef CONFIG_VMAP_STACK
        struct vm_struct          *stack_vm_area;
#endif

#ifdef CONFIG_THREAD_INFO_IN_TASK
        /* A live task holds one reference: */
        refcount_t                stack_refcount;
#endif

```

```

#ifdef CONFIG_LIVEPATCH
    int patch_state;
#endif
#ifdef CONFIG_SECURITY
    /* Used by LSM modules for access restriction: */
    void                      *security;
#endif

#ifdef CONFIG_GCC_PLUGIN_STACKLEAK
    unsigned long             lowest_stack;
    unsigned long             prev_lowest_stack;
#endif

    /*
     * New fields for task_struct should be added above here, so that
     * they are included in the randomized portion of task_struct.
     */
    randomized_struct_fields_end

    /* CPU-specific state of this task: */
    struct thread_struct       thread;

    /*
     * WARNING: on x86, 'thread_struct' contains a variable-sized
     * structure. It *MUST* be at the end of 'task_struct'.
     *
     * Do not put anything below here!
     */
};

```